

```
In [2]: 1 import warnings
2 warnings.filterwarnings('ignore')
3
4 import pandas as pd
5 import numpy as np
6 import seaborn as sns
7 import matplotlib.pyplot as plt
8 import librosa.display
9 import os
10
11 from sklearn.model_selection import train_test_split
```

```
In [3]: 1 path = '../input/cremad/AudioWAV/'
2 audio_path = []
3 audio_emotion = []
```

```
In [4]: 1 # collects all the audio filename in the variable 'path'
2 directory_path = os.listdir(path)
```

```
In [5]: 1 for audio in directory_path:
2     audio_path.append(path + audio)
3     emotion = audio.split('_')
4     if emotion[2] == 'SAD':
5         audio_emotion.append("sad")
6     elif emotion[2] == 'ANG':
7         audio_emotion.append("angry")
8     elif emotion[2] == 'DIS':
9         audio_emotion.append("disgust")
10    elif emotion[2] == 'NEU':
11        audio_emotion.append("neutral")
12    elif emotion[2] == 'HAP':
13        audio_emotion.append("happy")
14    elif emotion[2] == 'FEA':
15        audio_emotion.append("fear")
16    else:
17        audio_emotion.append("unknown")
```


```
In [6]: 1 emotion_dataset = pd.DataFrame(audio_emotion, columns=['Emotions'])
2 audio_path_dataset = pd.DataFrame(audio_path, columns=['Path'])
3 dataset = pd.concat([audio_path_dataset, emotion_dataset], axis= 1)
4 #print(Len(dataset))
5 print(dataset.head())
6 # print(dataset['File Path'][55])
```

	Path	Emotions
0	../input/cremad/AudioWAV/1028_TSI_DIS_XX.wav	disgust
1	../input/cremad/AudioWAV/1075_IEO_HAP_LO.wav	happy
2	../input/cremad/AudioWAV/1084_ITS_HAP_XX.wav	happy
3	../input/cremad/AudioWAV/1067_IWW_DIS_XX.wav	disgust
4	../input/cremad/AudioWAV/1066_TIE_DIS_XX.wav	disgust

## Visualization

```
In [7]: 1 # counting audio categorized by emotions
2 plt.figure(figsize=(6,6), dpi=80)
3 plt.title("Emotion Count", size=16)
4 plt.xlabel('Emotions', size = 12)
5 plt.ylabel('Count', size = 12)
6 sns.histplot(dataset.Emotions, color='#F19C0E')
7 #plt.show()
```

Out[7]: <AxesSubplot:title={'center':'Emotion Count'}, xlabel='Emotions', ylabel='Count'>



## Showing spectrogram and waveplot

```
In [8]: 1 emotion_sad = dataset[dataset['Emotions']=='sad']['Path']
2 print(type(emotion_sad))
```


<class 'pandas.core.series.Series'>

```
In [9]: 1 #choosing a file to plot wave and spectrogram
2 #print(emotion_sad.values[65])
3 data_path = emotion_sad.values[542]
4 data, sampling_rate = librosa.load(data_path)
```

## Waveplot

```
In [10]: 1 plt.figure(figsize=(10,6))
2 plt.title("Waveplot for a particular audio representing SAD emotion", size=16)
3 librosa.display.waveplot(data, sr=sampling_rate)
4 #plt.show()
```

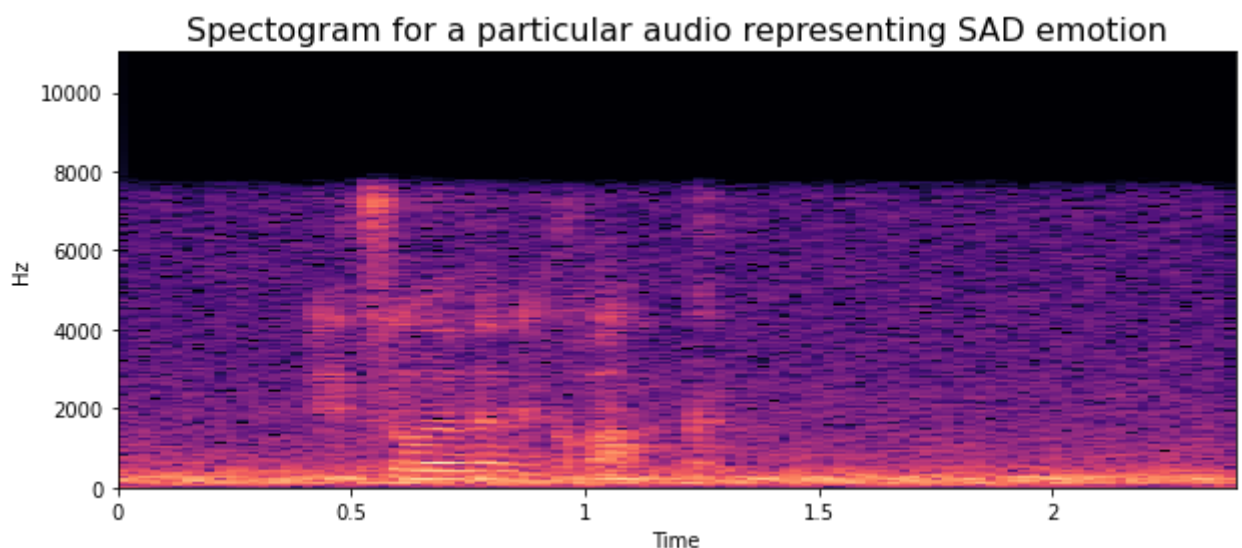
Out[10]: <matplotlib.collections.PolyCollection at 0x7f581b1dc1d0>



## Spectrogram

```
In [11]: 1 plt.figure(figsize=(10,4))
2 plt.title("Spectrogram for a particular audio representing SAD emotion", size=16)
3 D = librosa.stft(data)
4 S_db = librosa.amplitude_to_db(np.abs(D), ref=np.max)
5 librosa.display.specshow(S_db, sr = sampling_rate, x_axis='time', y_axis='hz')
6 #plt.show()
```

Out[11]: <matplotlib.collections.QuadMesh at 0x7f581bd23910>




# Augmentation (Noise Injection)

In [12]:

```
1 # for audio processing accuracy
2 # add noise to audio and check how the waveplot changes
3 # also the observing the change in audio quality
4
5 ## Augmentation (Noise Injection)
6 noise_amp = 0.035*np.random.uniform()*np.amax(data)
7 audio_injected_data = data + noise_amp*np.random.normal(size=data.shape[0])
8
9 # waveplot view after noise injection:
10 plt.figure(figsize=(10,6))
11 plt.title("Waveplot for a particular audio representing SAD emotion after noise inj
12 librosa.display.waveplot(audio_injected_data, sr=sampling_rate)
13 plt.show()
```

Waveplot for a particular audio representing SAD emotion after noise injection



## Feature Extraction

### Creating a DF with extracted Features

In [13]:

```
1 X, Y = [], []
2 print("Feature processing...")
3
4 for path, emo, index in zip(dataset.Path, dataset.Emotions, range(len(dataset))):
5     value, sample = librosa.load(path)
6     # noise injection
7     noise_amp = 0.035 * np.random.uniform() * np.amax(value)
8     value = value + noise_amp * np.random.normal(size=value.shape[0])
9     # mfcc
10    mfcc = librosa.feature.mfcc(y=value, sr= sample, n_mfcc=13, n_fft=200, hop_length=
11    mfcc = np.ravel(mfcc.T)
12    # mel
13    mel = librosa.feature.melspectrogram(y=value, sr=sample, hop_length = 256, n_fft=
14    mel = librosa.power_to_db(mel ** 2)
15    mel = np.ravel(mel).T
16    result = np.array([])
17    result = np.hstack((result, mfcc, mel))
18    #print(result)
19    result = np.array(result)
20    X.append(result)
21    Y.append(emo)
```

Feature processing...

In [14]:

```
1 # print(X)
2 # print(Y)
3 extracted_audio_df = pd.DataFrame(X)
4 extracted_audio_df["emotion_of_audio"] = Y
5 print(extracted_audio_df.shape)
6 print(extracted_audio_df.tail(10))
7 extracted_audio_df = extracted_audio_df.fillna(0)
8 #print(extracted_audio_df.isna().any())
```

(7442, 30457)

	0	1	2	3	4	5	\		
7432	-559.200843	-73.280942	-9.686211	-7.283784	18.869909	-8.773296			
7433	-512.106808	-58.911652	-10.667038	16.432381	20.050270	-6.125210			
7434	-512.564702	-78.250346	0.999713	10.071348	21.047835	1.798151			
7435	-687.661885	-91.652409	4.695701	28.066455	10.075825	-3.847991			
7436	-534.066959	-76.568546	-8.396436	-6.301705	1.900523	0.690286			
7437	-603.322967	-92.506487	10.069625	34.567998	-0.123639	-0.325672			
7438	-501.625926	-97.038078	-23.161478	9.334335	7.236941	9.445802			
7439	-544.170771	-91.829970	3.046557	12.402814	12.947438	13.671182			
7440	-792.501669	-46.832100	8.801988	33.063864	12.468825	-16.028599			
7441	-745.219608	-49.169030	12.641756	29.362595	-6.108275	-4.888654			
	6	7	8	9	...	30447	30448	30449	\
7432	-26.851384	0.918844	4.587700	-13.149765	...	NaN	NaN	NaN	
7433	-12.220169	-2.864142	-3.623588	3.218214	...	NaN	NaN	NaN	
7434	-30.792923	-18.055478	6.103335	3.545619	...	NaN	NaN	NaN	
7435	2.544490	-0.391296	-9.519397	-0.084294	...	NaN	NaN	NaN	
7436	-16.713967	-2.741473	4.235537	3.132909	...	NaN	NaN	NaN	
7437	14.165165	-9.085516	-13.073067	-3.160416	...	NaN	NaN	NaN	
7438	-8.530354	2.094583	17.751684	-14.345202	...	NaN	NaN	NaN	
7439	6.281240	-24.027286	-17.749574	13.954773	...	NaN	NaN	NaN	
7440	5.045405	2.159563	3.349647	-4.139884	...	NaN	NaN	NaN	
7441	3.836113	10.753297	6.675437	-12.665519	...	NaN	NaN	NaN	

	30450	30451	30452	30453	30454	30455	emotion_of_audio
7432	NaN	NaN	NaN	NaN	NaN	NaN	angry
7433	NaN	NaN	NaN	NaN	NaN	NaN	happy
7434	NaN	NaN	NaN	NaN	NaN	NaN	angry
7435	NaN	NaN	NaN	NaN	NaN	NaN	sad
7436	NaN	NaN	NaN	NaN	NaN	NaN	angry
7437	NaN	NaN	NaN	NaN	NaN	NaN	angry
7438	NaN	NaN	NaN	NaN	NaN	NaN	angry
7439	NaN	NaN	NaN	NaN	NaN	NaN	angry
7440	NaN	NaN	NaN	NaN	NaN	NaN	sad
7441	NaN	NaN	NaN	NaN	NaN	NaN	sad

[10 rows x 30457 columns]

## Training

In [15]:

```
1 # preparing to train
2 X = extracted_audio_df.drop(labels='emotion_of_audio', axis= 1)
3 Y = extracted_audio_df['emotion_of_audio']
4
5 x_train, x_test, y_train, y_test = train_test_split(np.array(X), Y, test_size=0.2)
```

```
In [16]: 1 from sklearn.preprocessing import StandardScaler
2 from sklearn.neural_network import MLPClassifier
3 from sklearn.metrics import accuracy_score
4
5
6 scaler = StandardScaler()
7 x_train = scaler.fit_transform(x_train)
8 x_test = scaler.transform(x_test)
```

## Model Creation and Fitting

```
In [17]: 1 mlp_model = MLPClassifier(activation='relu',
2                               solver='sgd',
3                               hidden_layer_sizes=100,
4                               alpha=0.839903176695813,
5                               batch_size=150,
6                               learning_rate='adaptive',
7                               max_iter=100000)
8 # Fit mlp model
9 mlp_model.fit(x_train,y_train)
```

```
Out[17]: MLPClassifier(alpha=0.839903176695813, batch_size=150, hidden_layer_sizes=100,
learning_rate='adaptive', max_iter=100000, solver='sgd')
```

## Accuracy Calculation

```
In [24]: 1 y_pred = mlp_model.predict(x_test)
2 accuracy=accuracy_score(y_true=y_test, y_pred=y_pred)
3
4 # the accuracy didn't turn out to be that good :(
5 print("\nModel:{} Accuracy: {:.2f}%".
6       format(type(mlp_model).__name__ , accuracy*100))
```

```
Model:MLPClassifier Accuracy: 39.29%
```

## Prediction Verification

```
In [25]: 1 # the prediction made by the model:
2 print("The Prediction Made By Model: ")
3 print("<<<=====>>>")
4 df = pd.DataFrame({'Actual': y_test, 'Predict': y_pred})
5 print(df.head())
```

```
The Prediction Made By Model:
```

```
<<<=====>>>
```

	Actual	Predict
7163	disgust	angry
5450	happy	neutral
3992	happy	happy
2441	fear	angry
2390	neutral	neutral

